

Project 2: UNIX Shell Programming & Linux Kernel Module for Task Information

1 UNIX Shell Programming

In this task, i will implement a UNIX Shell with following parts:

- (1) Creating the child process and executing the command in the child
- (2) Providing a history feature
- (3) Adding support of input and output redirection
- (4) Allowing the parent and child processes to communicate via a pipe

Design: My key designs is as follows:

- To parse the input command, i use **flatten()** function to remove extra space, then i use **tokenize()** function to transform string into array.
- After parsing, i check whether it is **!!** or **exit**, and give the result.
- Then i use **fork()** function to create the child process. The instruction will be execute in the child process.
- Then i use **concurrentFlag** function to check whether it asks for concurrent execution. If true, the parent process will not execute **wait(NULL)**. If false, the parent process will wait until the child process is done.
- In the child process, then i check whether the instruction needs a pipe. If true, **pipe(pipe_fd)** will be called to generate a pipe. It will be used to carry out data communication.
- If the instruction need to be redirected, **dup2()** function is used to redirect the input/output.
- As for the concrete instruction such as **ls**, i use **execvp()** to execute it.

The code of simple-shell.c is shown as follows.

```
1 #include <stdio.h>
2 #include <unistd.h>
3 #include <stdlib.h>
4 #include <string.h>
5 #include <sys/wait.h>
6 #include <fcntl.h>
7
8 #define MAX_LINE          80 /* 80 chars per line, per command */
9 #define bool              int /* simulate boolean... */
10 #define true              1
11 #define false             0
12
```

```
13 void flatten(char *a, char *b) { // remove extra space, \n, \t
14     int length = strlen(a);
15     int current_location = 0;
16     bool blank_end = false;
17     for (int i = 0; i < length; i++) {
18         if (a[i] == ' ' || a[i] == '\t' || a[i] == '\n') {
19             if (blank_end) {
20                 b[current_location++] = ' ';
21             }
22             blank_end = false;
23         } else {
24             b[current_location++] = a[i];
25             blank_end = true;
26         }
27     }
28     if (b[current_location - 1] == ' ') {
29         b[current_location - 1] = 0;
30     }
31     return;
32 }
33
34 int tokenize(char *a, char **b) { //tokenize flatten string
35     int num = 0; // token num
36     int current_location = 0; // location inside string
37     int length = strlen(a);
38     if (length == 0) return 0;
39     for(int i = 0; i < length; i++) {
40         if (a[i] == ' ') {
41             num++;
42             current_location = 0;
43         } else {
44             b[num][current_location++] = a[i];
45         }
46     }
47     return num + 1;
48 }
49
50 bool concurrentFlag(char *a) {
51     int length = strlen(a);
52     if (length == 0) return false;
53     return (a[length - 1] == '&');
54 }
55
56 int main(void) {
57     char *args[MAX_LINE/2 + 1];    /* command line (of 80) has max of 40 arguments
58                                     */
59     char *raw_string;
60     char *flatten_string;
61     char *last_command;
62     int should_run = 1;
63     bool have_last_command = false;
```

```
63     raw_string = (char*) malloc(MAX_LINE * sizeof(char));
64     flatten_string = (char*) malloc(MAX_LINE * sizeof(char));
65     last_command = (char*) malloc(MAX_LINE * sizeof(char));
66     while (should_run){
67         if (concurrentFlag(flatten_string) == false) wait(NULL);
68         printf("\033[42;37msh>\033[2;7;0m"); // pretty print ~
69         fflush(stdout);
70         memset(raw_string, 0, sizeof(raw_string));
71         memset(flatten_string, 0, sizeof(raw_string));
72         fflush(stdout);
73         fgets(raw_string, MAX_LINE, stdin);
74         flatten(raw_string, flatten_string); // remove space \t \n from
75         raw_string
76
77         if (strcmp(flatten_string, "exit") == 0) {
78             should_run = 0;
79             break;
80         }
81         if (strcmp(flatten_string, "!!") == 0) {
82             if (have_last_command) {
83                 strcpy(flatten_string, last_command);
84                 fprintf(stdout, "%s\n", flatten_string);
85             } else {
86                 fprintf(stderr, "No commands in history.\n");
87                 continue;
88             }
89         } else {
90             have_last_command = true;
91             strcpy(last_command, flatten_string);
92         }
93
94         pid_t pid;
95         pid = fork();
96         if (pid < 0) {
97             fprintf(stderr, "Failed when creating a child process.\n");
98             continue;
99         }
100         if (pid == 0) { // child
101             for (int i = 0; i < MAX_LINE/2 + 1; i++) {
102                 args[i] = (char*) malloc(MAX_LINE * sizeof(char));
103                 memset(args[i], 0, sizeof(args[i]));
104             }
105             int command_num = tokenize(flatten_string, args);
106             for (int i = command_num; i < MAX_LINE/2 + 1; i++) {
107                 free(args[i]);
108                 args[i] = NULL;
109             }
110
111             if (concurrentFlag(flatten_string)) {
112                 free(args[command_num - 1]);
113                 args[command_num - 1] = NULL;
```

```
113         command_num = command_num - 1;
114     }
115
116     int pipe_loc = -1;
117     for (int i = 0; i < command_num; i++) {
118         if (strcmp(args[i], "|") == 0) {
119             pipe_loc = i;
120             break;
121         }
122     }
123
124     if (pipe_loc != -1) { // if pipe is need
125         if (pipe_loc == 0 || pipe_loc == command_num - 1) {
126             fprintf(stderr, "Pipe command can't be conducted,
127                 expect more commands.\n");
128         } else {
129             int pipe_fd[2];
130             int pipe_flag = pipe(pipe_fd);
131             if (pipe_flag == -1) {
132                 fprintf(stderr, "Pipe can't be
133                     established.\n");
134             } else { // pipe is created and no error, just do
135                 it ~
136                 pid = fork();
137                 if (pid < 0) {
138                     fprintf(stderr, "Failed when
139                         creating a child process.\n");
140                     ;
141                 }
142                 if (pid == 0) { // child of child
143                     for (int i = pipe_loc; i <
144                         command_num; i++) {
145                         free(args[i]);
146                         args[i] = NULL;
147                     }
148
149                     close(pipe_fd[0]); // close read
150                     if (dup2(pipe_fd[1],
151                         STDOUT_FILENO) < 0) {
152                         fprintf(stderr, "Dup
153                             failed.\n");
154                     } else {
155                         execvp(args[0], args);
156                     }
157                     close(pipe_fd[1]); // close write
158
159                     for (int i = 0; i < pipe_loc; i
160                         ++){ // child of child
161                         memory free
162                         free(args[i]);
163                     }
164                 }
165             }
166         }
167     }
168 }
```

```
154         free(raw_string);
155         free(flatten_string);
156         free(last_command);
157         exit(0);
158     } else { // child
159         wait(NULL);
160
161         for (int i = 0; i <= pipe_loc; i
162             ++){
163             free(args[i]);
164         }
165         for (int i = pipe_loc + 1; i <
166             command_num; i++) { // move 2
167             nd command to 1st
168                 args[i - pipe_loc - 1] =
169                     args[i];
170         }
171         for (int i = command_num -
172             pipe_loc - 1; i < command_num
173             ; i++) {
174             args[i] = NULL;
175         }
176         command_num = command_num -
177             pipe_loc - 1;
178
179         close(pipe_fd[1]);
180         if (dup2(pipe_fd[0], STDIN_FILENO
181             ) < 0) {
182             fprintf(stderr, "Pipe dup
183                 failed.\n");
184         } else {
185             execvp(args[0], args);
186         }
187         close(pipe_fd[0]);
188     }
189 }
190
191 } else { // no pipe, consider redirect input/output
192     bool input_redirect_flag = false;
193     bool output_redirect_flag = false;
194     char *file_name;
195     file_name = (char*) malloc(MAX_LINE * sizeof(char));
196     if (command_num >= 3) {
197         if ((strcmp(args[command_num - 2], "<") == 0) ||
198             (strcmp(args[command_num - 2], ">") == 0)) {
199             if (strcmp(args[command_num - 2], "<") ==
200                 0) {
201                 input_redirect_flag = true;
202             } else {
203                 output_redirect_flag = true;
204             }
205         }
206     }
207 }
```

```
194         strcpy(file_name, args[command_num - 1]);
195
196         free(args[command_num - 1]);
197         free(args[command_num - 2]);
198         args[command_num - 1] = NULL;
199         args[command_num - 2] = NULL;
200         command_num = command_num - 2;
201     }
202 }
203 int file_fd;
204 if (input_redirect_flag == true) {
205     file_fd = open(file_name, O_RDONLY, 0644);
206     if (file_fd < 0) {
207         fprintf(stderr, "File doesn't exist.\n");
208     } else {
209         if (dup2(file_fd, STDIN_FILENO) < 0) {
210             fprintf(stderr, "Dup failed.\n");
211         } else {
212             execvp(args[0], args);
213             close(file_fd);
214         }
215     }
216 } else if (output_redirect_flag == true) {
217     file_fd = open(file_name, O_WRONLY | O_CREAT,
218         0644);
219     if (file_fd < 0) {
220         fprintf(stderr, "Can't create file.\n");
221     } else {
222         if (dup2(file_fd, STDOUT_FILENO) < 0) {
223             fprintf(stderr, "Dup failed.\n");
224         } else {
225             execvp(args[0], args);
226             close(file_fd);
227         }
228     }
229 } else {
230     execvp(args[0], args); // normal, no pipe, no
231                             // redirect
232 }
233 free(file_name);
234 }
235
236 for (int i = 0; i < command_num; i++) { // child memory free
237     free(args[i]);
238 }
239 free(raw_string);
240 free(flatten_string);
241 free(last_command);
242 exit(0);
243 } else { // father
244     if (concurrentFlag(flatten_string) == false) wait(NULL);
```

```

243         }
244     }
245
246     free(raw_string);
247     free(flatten_string);
248     free(last_command);
249     return 0;
250 }

```

Since it is a simple program, i use the following command to compile it. (without Makefile)

```
1 gcc simple-shell.c -o shell
```

The result is shown as follows.

```

misaka@ubuntu:~/Documents/project2$ gcc simple-shell.c -o shell
misaka@ubuntu:~/Documents/project2$ ./shell
osh>!!
No commands in history.
osh>ls -a
. . . shell simple-shell simple-shell.c test_redirect
osh>!!
ls -a
. . . shell simple-shell simple-shell.c test_redirect
osh>ls -l > test_output
osh>sort < test_output
-rw-r--r-- 1 misaka misaka  0 Apr  6 22:36 test_output
-rw-r--r-- 1 misaka misaka 182 Apr  6 20:30 test_redirect
-rwxrwx-rw- 1 misaka misaka 6781 Apr  6 22:35 simple-shell.c
-rwxrwxr-x 1 misaka misaka 21848 Apr  6 22:34 simple-shell
-rwxrwxr-x 1 misaka misaka 21848 Apr  6 22:36 shell
total 60
osh>ls -a
. . . shell simple-shell simple-shell.c test_output test_redirect
osh>ls -a &
osh>. . . shell simple-shell simple-shell.c test_output test_redirect
ls
shell simple-shell simple-shell.c test_output test_redirect
osh>ls -l | sort
-rw-r--r-- 1 misaka misaka  182 Apr  6 20:30 test_redirect
-rw-r--r-- 1 misaka misaka  299 Apr  6 22:36 test_output
-rwxrwx-rw- 1 misaka misaka 6781 Apr  6 22:35 simple-shell.c
-rwxrwxr-x 1 misaka misaka 21848 Apr  6 22:34 simple-shell
-rwxrwxr-x 1 misaka misaka 21848 Apr  6 22:36 shell
total 64
osh>exit
misaka@ubuntu:~/Documents/project2$

```

图 1: Shell

2 Linux Kernel Module for Task Information

In this project, we are required to write a Linux kernel module that use the `/proc` file system for displaying a task's information based on its process identifier value `pid`.

Design: My key design is as follows:

- In `proc_write()`, i use `sscanf()` function to get the `pid` value from input.
- Memory allocation is different in the kernel. I use `kmalloc()` and `kfree()` to allocate memory.
- In `proc_read()`, i use `pid_task(find_vpid(...), ...)` function to read the PCB information with a `pid` number.
- When there is a error (not valid `pid` number), i will print the kernel information to inform the user.

The code of pid.c is shown as follows.

```
1  #include <linux/init.h>
2  #include <linux/slab.h>
3  #include <linux/sched.h>
4  #include <linux/module.h>
5  #include <linux/kernel.h>
6  #include <linux/proc_fs.h>
7  #include <linux/vmalloc.h>
8  #include <asm/uaccess.h>
9
10 #define BUFFER_SIZE 128
11 #define PROC_NAME "pid"
12
13 /* the current pid */
14 static long l_pid;
15
16 static ssize_t proc_read(struct file *file, char *buf, size_t count, loff_t *pos);
17 static ssize_t proc_write(struct file *file, const char __user *usr_buf, size_t count,
18     loff_t *pos);
19
20 static struct proc_ops proc_op = {
21     .proc_read = proc_read,
22     .proc_write = proc_write
23 };
24
25 static int proc_init(void)
26 {
27     proc_create(PROC_NAME, 0666, NULL, &proc_op);
28     printk(KERN_INFO "/proc/%s created\n", PROC_NAME);
29     return 0;
30 }
31
32 static void proc_exit(void)
33 {
34     remove_proc_entry(PROC_NAME, NULL);
35     printk( KERN_INFO "/proc/%s removed\n", PROC_NAME);
36 }
37
38 /**
39  * This function is called each time the /proc/pid is read.
40  *
41  * This function is called repeatedly until it returns 0, so
42  * there must be logic that ensures it ultimately returns 0
43  * once it has collected the data that is to go into the
44  * corresponding /proc file.
45  */
46 static ssize_t proc_read(struct file *file, char __user *usr_buf, size_t count, loff_t *
47     pos)
48 {
49     int rv = 0;
50     char buffer[BUFFER_SIZE];
```



```
49     static int completed = 0;
50     struct task_struct *tsk = NULL;
51
52     if (completed) {
53         completed = 0;
54         return 0;
55     }
56
57     tsk = pid_task(find_vpid(l_pid), PIDTYPE_PID);
58     if (tsk == NULL) {
59         printk(KERN_INFO "Not a valid pid %ld\n", l_pid);
60         return 0;
61     }
62     completed = 1;
63     rv = sprintf(buffer, "command = [%s] pid = [%ld] state = [%ld]\n", tsk->comm,
64                  l_pid, tsk->state);
65     if (copy_to_user(usr_buf, buffer, rv)) {
66         rv = -1;
67     }
68
69     return rv;
70 }
71
72 /**
73  * This function is called each time we write to the /proc file system.
74  */
75 static ssize_t proc_write(struct file *file, const char __user *usr_buf, size_t count,
76                          loff_t *pos)
77 {
78     char *k_mem;
79
80     k_mem = kmalloc(count, GFP_KERNEL);
81
82     if (copy_from_user(k_mem, usr_buf, count)) {
83         printk(KERN_INFO "Error copying from user\n");
84         return -1;
85     }
86
87     sscanf(k_mem, "%ld", &l_pid);
88
89     kfree(k_mem);
90
91     return count;
92 }
93
94 /** Macros for registering module entry and exit points. */
95 module_init( proc_init );
96 module_exit( proc_exit );
97
98 MODULE_LICENSE("GPL");
99 MODULE_DESCRIPTION("Pid");
```

```
98 MODULE_AUTHOR("MisakaCenter");
```

The Makefile for this task is written as follows.

```
1 obj-m := pid.o
2 all:
3     make -C /usr/src/linux-5.11.3/ M=$(shell pwd) modules
4 clean:
5     make -C /usr/src/linux-5.11.3/ M=$(shell pwd) clean
```

The result is shown as follows.

```
misaka@ubuntu:~/Documents/project2/pid$ make
make -C /usr/src/linux-5.11.3/ M=/home/misaka/Documents/project2/pid modules
make[1]: Entering directory '/usr/src/linux-5.11.3'
  CC [M]  /home/misaka/Documents/project2/pid/pid.o
  MODPOST /home/misaka/Documents/project2/pid/Module.symvers
  CC [M]  /home/misaka/Documents/project2/pid/pid.mod.o
  LD [M]  /home/misaka/Documents/project2/pid/pid.ko
make[1]: Leaving directory '/usr/src/linux-5.11.3'
misaka@ubuntu:~/Documents/project2/pid$ sudo dmesg -C
misaka@ubuntu:~/Documents/project2/pid$ sudo insmod pid.ko
misaka@ubuntu:~/Documents/project2/pid$ sudo dmesg
[13443.032003] /proc/pid created
misaka@ubuntu:~/Documents/project2/pid$ echo "1" > /proc/pid
misaka@ubuntu:~/Documents/project2/pid$ cat /proc/pid
command = [systemd] pid = [1] state = [1]
misaka@ubuntu:~/Documents/project2/pid$ echo "233" > /proc/pid
misaka@ubuntu:~/Documents/project2/pid$ cat /proc/pid
command = [kthreadd] pid = [2] state = [1]
misaka@ubuntu:~/Documents/project2/pid$ sudo rmmod pid
misaka@ubuntu:~/Documents/project2/pid$ sudo dmesg
[13443.032003] /proc/pid created
[13485.801628] Not a valid pid 233
[13540.301943] /proc/pid removed
```

图 2: Task information