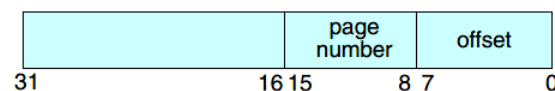


Project 8: Designing a Virtual Memory Manager

This project consists of writing a program that translates logical to physical addresses for a virtual address space of size $2^{16} = 65536$ bytes. The program will read from a file containing logical address and, using a TLB and a page table, will translate each logical address to its corresponding physical address and output the value of the byte stored at the translated physical address. It requires us to use simulation to understand the steps involved in translating logical address to physical address, and this will include resolving page faults using demand paging, managing a TLB, and implementing a page-replacement algorithm.

The program will read a file containing several 32-bit integer numbers that represent logical addresses. However, the 16-bit addresses is the only thing that needs to be concerned, so we must mask the rightmost 16 bits of each logical addresses. These 16 bits are divided into an 8-bit page number and an 8-bit page offset. Hence, the addresses are structured as shown as:



Other specifics includes the following:

- 2^8 entries in the page table;
- Page size of 2^8 bytes;
- 16 entries in the TLB;
- Frame size of 2^8 bytes;
- 256 frames;
- Physical memory of 65536 bytes ($256 \text{ frames} \times 256\text{-byte frame size}$).

The program is to output the following values:

- The logical address being translated (the integer value being read from **addresses.txt**).
- The corresponding physical address (what your program translates the logical address to).
- The signed byte value stored in physical memory at the translated physical address.

We also provide the file **correct.txt**, which contains the correct output values for the file **addresses.txt**. You should use this file to determine if the program is correctly translating logical to physical addresses.

After completion, the program is to report the following statistics:

- Page-fault rate - The percentage of address references that resulted in page faults.
- TLB hit rate - The percentage of address references that were resolved in the TLB.

Since the logical addresses in `addresses.txt` were generated randomly and do not reflect any memory access locality, do not expect to have a high TLB hit rate.

Then use a smaller physical address space with 128 page frames rather than 256. This change will require modifying your program so that it keeps track of free page frames as well as implementing a page-replacement policy using either FIFO or LRU (Section 10.4) to resolve page faults when there is no free memory.

Design: My design for this task is:

- All the replacement (memory and TLB) is implemented with LRU algorithm.
- `swap_page_in()` function will first search for the empty frame. If there is no empty frame, LRU algorithm will be used to find a victim to swap. The victim will be deleted from TLB using `delete_TLB()` function.
- For translating a virtual address to a physical one, the program will first look up TLB using `get_frame_TLB()` function. If TLB miss, it will look up `page_table` with `page_table_inmem`. If page table miss, it will swap the page in from `backing_store` using `swap_page_in()` function.

The implementation of the virtual memory manager (`vmm.c`) is shown as follows.

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <string.h>
4
5  #define PAGE_NUM 256
6  #define PAGE_SIZE 256
7  #define FRAME_NUM 256
8  #define FRAME_SIZE 256
9  #define TLB_SIZE 16
10
11 int get_frame_TLB(int page_number);
12 void TLB_update(int page_number, int frame_number);
13 void delete_TLB(int page_number, int frame_number);
14
15 char memory[FRAME_NUM * FRAME_SIZE];
16
17 FILE *backing_store;
18 int total = 0, page_fault = 0, tlb_hit = 0;
19
20 int page_table[PAGE_NUM], page_table_inmem[PAGE_NUM];
21 int frame_count[FRAME_NUM];
22
23 void swap_page_out(int frame_number) {
24     int page_number = -1;
25
26     for (int i = 0; i < PAGE_NUM; i++) {
27         if (page_table_inmem[i] > 0 && page_table[i] == frame_number) {
28             page_number = i;
29             break;
30         }
31     }
32     if (page_number == -1) {
33         for (int i = 0; i < PAGE_NUM; i++) {
34
```

```
35     printf("%d, %d \n", page_table_inmem[i], page_table[i]);
36 }
37     fprintf(stderr, "[Err] Unexpected Error!\n");
38
39     exit(1);
40 }
41 page_table_inmem[page_number] = 0;
42
43 //TLB delete
44 delete_TLB(page_number, frame_number);
45
46 return;
47 }
48
49 int swap_page_in(int page_number) {
50     char buf[FRAME_SIZE];
51     fseek(backing_store, page_number * FRAME_SIZE, 0);
52     fread(buf, sizeof(char), FRAME_SIZE, backing_store);
53
54     int target = -1;
55     for (int i = 0; i < FRAME_NUM; i++) {
56         if (frame_count[i] == 0) {
57             target = i;
58             break;
59         }
60     }
61     if (target == -1) {
62         for (int i = 0; i < FRAME_NUM; i++) {
63             if (frame_count[i] == FRAME_NUM) {
64                 target = i;
65                 break;
66             }
67         }
68         swap_page_out(target);
69     }
70     for (int i = 0; i < FRAME_SIZE; ++i) memory[target * FRAME_SIZE + i] = buf[i];
71     for (int i = 0; i < FRAME_NUM; ++i) if (frame_count[i] > 0) frame_count[i]++;
72     frame_count[target] = 1;
73     return target;
74 }
75
76
77 int get_val(int frame_number, int offset) {
78     int val = (int)memory[frame_number * FRAME_SIZE + offset];
79
80     for (int i = 0; i < FRAME_NUM; ++i) {
81         if (frame_count[i] != 0 && frame_count[i] < frame_count[frame_number])
82             frame_count[i]++;
83     }
84     frame_count[frame_number] = 1;
85 }
```

```
86     return val;
87 }
88
89 int get_frame(int page_number) {
90
91     // TLB
92     int tlb_frame = get_frame_TLB(page_number);
93     if (tlb_frame != -1) return tlb_frame;
94
95     if (page_table_inmem[page_number] > 0) {
96
97         // TLB
98         TLB_update(page_number, page_table[page_number]);
99
100         return page_table[page_number];
101     } else {
102         page_fault++;
103         page_table[page_number] = swap_page_in(page_number);
104         page_table_inmem[page_number] = 1;
105
106         // TLB
107         TLB_update(page_number, page_table[page_number]);
108         return page_table[page_number];
109     }
110     return -1;
111 }
112
113
114 // TLB
115
116 int TLB_list[TLB_SIZE][2]; // 0 -> page; 1 -> frame
117 int TLB_count[TLB_SIZE];
118
119 int get_frame_TLB(int page_number) {
120     int frame = -1, loc = -1;
121     for (int i = 0; i < TLB_SIZE; ++i)
122         if (TLB_count[i] != 0 && TLB_list[i][0] == page_number) {
123             frame = TLB_list[i][1];
124             loc = i;
125             break;
126         }
127
128     if (frame == -1) return -1; // TLB miss
129
130     tlb_hit++;
131     for (int i = 0; i < TLB_SIZE; i++)
132         if (TLB_count[i] != 0 && TLB_count[i] < TLB_count[loc]) {
133             TLB_count[i]++;
134         }
135     TLB_count[loc] = 1;
136 }
```

```
137     return frame;
138 }
139
140 void TLB_update(int page_number, int frame_number) {
141     int loc = -1;
142     for (int i = 0; i < TLB_SIZE; i++)
143         if(TLB_count[i] == 0) {
144             loc = i;
145             break;
146         }
147     if (loc == -1) {
148         for (int i = 0; i < TLB_SIZE; i++)
149             if(TLB_count[i] == TLB_SIZE) {
150                 loc = i;
151                 break;
152             }
153     }
154
155     for (int i = 0; i < TLB_SIZE; i++)
156         if (TLB_count[i] > 0) TLB_count[i] += 1;
157
158     TLB_count[loc] = 1;
159     TLB_list[loc][0] = page_number;
160     TLB_list[loc][1] = frame_number;
161     return;
162 }
163
164 void delete_TLB(int page_number, int frame_number) {
165     int loc = -1;
166     for (int i = 0; i < TLB_SIZE; ++ i)
167         if(TLB_count[i] == 0) {
168             loc = i;
169             break;
170         }
171     if (loc == -1) return; // not find, no need to delete
172
173     for (int i = 0; i < TLB_SIZE; ++ i)
174         if (TLB_count[i] > TLB_count[loc]) TLB_count[i] -= 1;
175
176     TLB_count[loc] = 0;
177 }
178
179
180 int main(int argc, char *argv[]) {
181     if (argc != 2) {
182         fprintf(stderr, "[Error] Unexpected input!\n");
183         return 1;
184     }
185
186     backing_store = fopen("BACKING_STORE.bin", "r");
187
```

```

188     for (int i = 0; i < PAGE_NUM; i++) {
189         page_table[i] = 0;
190         page_table_inmem[i] = 0;
191     }
192     for (int i = 0; i < FRAME_NUM; i++) {
193         frame_count[i] = 0;
194     }
195     for (int i = 0; i < TLB_SIZE; i++) {
196         TLB_list[i][0] = 0;
197         TLB_list[i][1] = 0;
198         TLB_count[i] = 0;
199     }
200
201     FILE *stream_in = fopen(argv[1], "r");
202     FILE *stream_out = fopen("output.txt", "w");
203
204     int virtual_addr;
205     int page_number, offset, frame_number, val;
206
207     while (~fscanf(stream_in, "%d", &virtual_addr)) {
208         total += 1;
209
210         virtual_addr = virtual_addr & 0x0000ffff;
211         page_number = (virtual_addr >> 8) & 0x000000ff;
212         offset = virtual_addr & 0x000000ff;
213
214         frame_number = get_frame(page_number);
215         val = get_val(frame_number, offset);
216
217         fprintf(stream_out, "Virtual address: %d Physical address: %d Value: %d\n",
218             virtual_addr, (frame_number << 8) + offset, val);
219     }
220     fprintf(stdout, "TLB hit rate: %.2f%%      Page fault rate: %.2f%%\n", 100.0 *
221         tlb_hit / total, 100.0 * page_fault / total);
222
223     fclose(stream_in);
224     fclose(stream_out);
225
226     return 0;
227 }

```

To check the result, i also implement a checker program (`judge.c`), the main function of it is to compare the value of the answer and the output. If the output is correct, this program will print **All correct!**, otherwise, it will print **Wrong answer!**. The code of `judge.c` is shown as follow.

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <string.h>
4
5  int main() {
6      FILE *stream_ans = fopen("output.txt", "r");
7      FILE *stream_std = fopen("correct.txt", "r");

```

```
8
9  int a, b, c;
10 int ac = 0; // 0 => ac; 1 => wrong
11
12 while(~fscanf(stream_ans, "Virtual address: %d Physical address: %d Value: %d\n", &a,
13     &b, &c)) {
14     int d, e, f;
15     if (fscanf(stream_std, "Virtual address: %d Physical address: %d Value: %d\n", &d
16         , &e, &f) == EOF) {
17         printf("File length not match!\n");
18         ac = 1;
19         break;
20     }
21     if (c != f) {
22         printf("Wrong answer!\n");
23         ac = 1;
24         break;
25     }
26 }
27
28 if (ac == 0) printf("All correct!\n");
29
30 fclose(stream_ans);
31 fclose(stream_std);
32
33 return 0;
34 }
```

Makefile for this task is shown as below:

```
1 CC=gcc
2 CFLAGS=-Wall
3
4 all: vmm.o judge.o
5     $(CC) $(CFLAGS) -o vmm vmm.o
6     $(CC) $(CFLAGS) -o judge judge.o
7
8 vmm.o: vmm.c
9     $(CC) $(CFLAGS) -c vmm.c
10
11 judge.o: judge.c
12     $(CC) $(CFLAGS) -c judge.c
13
14 clean:
15     rm -rf *.o
16     rm -rf vmm
17     rm -rf judge
```

When the **FRAME_NUM** is 128, the execution result of the virtual memory manager is shown as follows:

```
misaka@MS-BVZPMBEQIPCD:/mnt/c/Projects/OS_Project/Project8/project8$ make all
gcc -Wall -c vmm.c
gcc -Wall -c judge.c
gcc -Wall -o vmm vmm.o
gcc -Wall -o judge judge.o
misaka@MS-BVZPMBEQIPCD:/mnt/c/Projects/OS_Project/Project8/project8$ ./vmm addresses.txt
TLB hit rate: 5.50%      Page fault rate: 53.90%
misaka@MS-BVZPMBEQIPCD:/mnt/c/Projects/OS_Project/Project8/project8$ ./judge
All correct!
```

图 1: Designing a Virtual Memory Manager (frame number is 128)

When the **FRAME_NUM** is 256, the execution result of the virtual memory manager is shown as follows. The page fault rate is dropped.

```
misaka@MS-BVZPMBEQIPCD:/mnt/c/Projects/OS_Project/Project8/project8$ make all
gcc -Wall -c vmm.c
gcc -Wall -c judge.c
gcc -Wall -o vmm vmm.o
gcc -Wall -o judge judge.o
misaka@MS-BVZPMBEQIPCD:/mnt/c/Projects/OS_Project/Project8/project8$ ./vmm addresses.txt
TLB hit rate: 5.50%      Page fault rate: 24.40%
misaka@MS-BVZPMBEQIPCD:/mnt/c/Projects/OS_Project/Project8/project8$ ./judge
All correct!
```

图 2: Designing a Virtual Memory Manager (frame number is 256)